

Comprehensive Predictive Modeling of Ames House Prices

June 17, 2024

1 Comprehensive Predictive Modeling of Ames House Prices

1. Introduction (200 words) • Overview of Real Estate Price Prediction (100 words): Discuss the significance of predicting real estate prices, its impact on industries like marketing, business intelligence, and urban planning. • Scope and Objectives (50 words): Outline the aims of the research, including developing predictive models and evaluating their effectiveness. • Structure of the Report (50 words): Briefly describe the structure of the report, summarizing the main sections.

2. Exploratory Data Analysis

2.1 Dataset Description

- Description of the Dataset (100 words): Explain the key attributes (e.g., zoning, lot size, building characteristics, neighborhood factors).
- Source (25 words): Cite the dataset source (e.g., Kaggle, UCI Machine Learning Repository).

Import libraries

```
[ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import scipy.stats as stats
from scipy.stats import skew

%matplotlib inline
sns.set_style('darkgrid')
```

Load data

```
[ ]: dataset = pd.read_csv("house-price-data-apr2024.csv")
dataset.describe()
```

```
[ ]:      LotFrontage      LotArea  OverallQual  OverallCond  MasVnrArea  \
count    1201.000000    1460.000000    1460.000000    1460.000000    1452.000000
mean         70.049958    10516.828082         6.099315         5.575342     103.685262
std         24.284752     9981.264932         1.382997         1.112799     181.066207
min         21.000000     1300.000000         1.000000         1.000000         0.000000
```

25%	59.000000	7553.500000	5.000000	5.000000	0.000000
50%	69.000000	9478.500000	6.000000	5.000000	0.000000
75%	80.000000	11601.500000	7.000000	6.000000	166.000000
max	313.000000	215245.000000	10.000000	9.000000	1600.000000

	BsmtFinSF1	BsmtFinSF2	BsmtUnfSF	TotalBsmtSF	1stFlrSF	...	\
count	1460.000000	1460.000000	1460.000000	1460.000000	1460.000000	...	
mean	443.639726	46.549315	567.240411	1057.429452	1162.626712	...	
std	456.098091	161.319273	441.866955	438.705324	386.587738	...	
min	0.000000	0.000000	0.000000	0.000000	334.000000	...	
25%	0.000000	0.000000	223.000000	795.750000	882.000000	...	
50%	383.500000	0.000000	477.500000	991.500000	1087.000000	...	
75%	712.250000	0.000000	808.000000	1298.250000	1391.250000	...	
max	5644.000000	1474.000000	2336.000000	6110.000000	4692.000000	...	

	GarageCars	GarageArea	WoodDeckSF	OpenPorchSF	EnclosedPorch	...	\
count	1460.000000	1460.000000	1460.000000	1460.000000	1460.000000	...	
mean	1.767123	472.980137	94.244521	46.660274	21.954110	...	
std	0.747315	213.804841	125.338794	66.256028	61.119149	...	
min	0.000000	0.000000	0.000000	0.000000	0.000000	...	
25%	1.000000	334.500000	0.000000	0.000000	0.000000	...	
50%	2.000000	480.000000	0.000000	25.000000	0.000000	...	
75%	2.000000	576.000000	168.000000	68.000000	0.000000	...	
max	4.000000	1418.000000	857.000000	547.000000	552.000000	...	

	3SsnPorch	ScreenPorch	PoolArea	MiscVal	SalePrice
count	1460.000000	1460.000000	1460.000000	1460.000000	1460.000000
mean	3.409589	15.060959	2.758904	43.489041	180921.195890
std	29.317331	55.757415	40.177307	496.123024	79442.502883
min	0.000000	0.000000	0.000000	0.000000	34900.000000
25%	0.000000	0.000000	0.000000	0.000000	129975.000000
50%	0.000000	0.000000	0.000000	0.000000	163000.000000
75%	0.000000	0.000000	0.000000	0.000000	214000.000000
max	508.000000	480.000000	738.000000	15500.000000	755000.000000

[8 rows x 32 columns]

2.1 Feature Distributions Relative to target

A comprehensive initial assessment of how each of the 46 features is distributed relative to the target variable `SalePrice` is visualized.

```
[ ]: # Assuming 'dataset' is your DataFrame
df = dataset.copy() # Make a copy to avoid modifying the original dataset

# List of features to visualize
features = df.columns.tolist()
features.remove('SalePrice')
```

```

target = 'SalePrice'

# Number of rows and columns for subplots
num_rows = 8
num_cols = 6

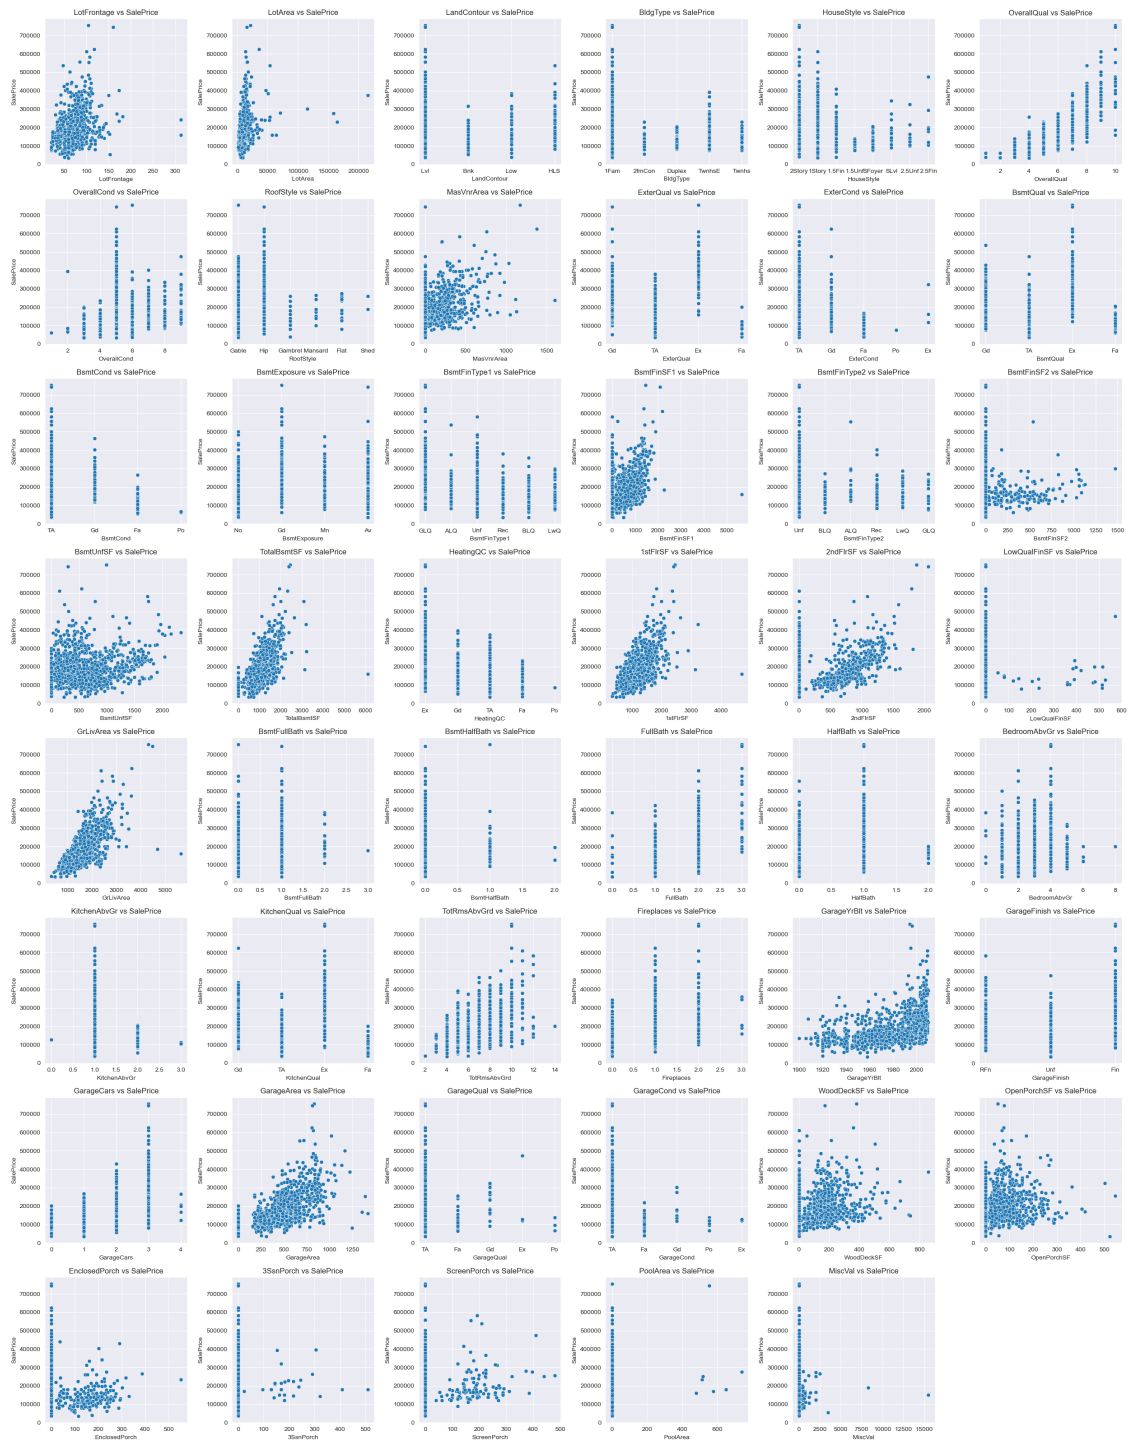
# Create subplots
fig, axes = plt.subplots(num_rows, num_cols, figsize=(25, 4*num_rows))
axes = axes.flatten()

# Plot each feature and calculate skewness
for i, feature in enumerate(features):
    # Plotting scatterplot
    sns.scatterplot(x=df[feature], y=df[target], ax=axes[i])
    axes[i].set_title(f'{feature} vs {target}')
    axes[i].set_xlabel(feature)
    axes[i].set_ylabel(target)

# Remove any unused subplots
for j in range(i + 1, len(axes)):
    fig.delaxes(axes[j])

plt.tight_layout()
plt.show()

```



2.2 Sale Price

```
[ ]: saleprice = dataset["SalePrice"]
sns.displot(saleprice,kde = True)
```

```
plt.title("House Price Distribution")
plt.show()
```



From the plot the distribution of `SalePrice` data is apparently right skewed. Its Skewness and Kurtosis statistics are checked. This was expected as few people can afford very expensive houses.

```
[ ]: print(f"""Skewness: {saleprice.skew()}
Kurtosis: {saleprice.kurt()}""")
```

```
Skewness: 1.8828757597682129
Kurtosis: 6.536281860064529
```

2.3 Numerical Features

There are total 31 numerical features.

```
[ ]: numerical_data = df.select_dtypes(include="number")
numerical_list = numerical_data.columns.tolist()
numerical_list.remove("SalePrice")
```

```
print("Number of numerical variables " , len(numerical_list))
```

Number of numerical variables 31

Correlation coefficient with target variable


To have a first glance of which numeric variables have a high correlation with the **SalePrice**. The correlation coefficient is computed. All coefficient > 0.4 is initially defined as a relative strong linear relationship and visualized. However, it is clear that the multicollinearity is an issue. For example: the correlation between **GarageCars** and **GarageArea** is very high (0.89), while both have relatively high correlations with **SalePrice**.

Pearson Correlation Range	Strength of Correlation	Interpretation
0.7 to 1.0 (positive or negative)	Strong	Indicates a strong linear relationship where one variable tends to increase (or decrease) as the other variable increases.
0.5 to 0.7 (positive or negative)	Moderate to Strong	Suggests a meaningful linear relationship between the variables, though not as strong as the highest range.
0.3 to 0.5 (positive or negative)	Weak to Moderate	The relationship exists but is less pronounced and may not be as influential in linear models.
0 to 0.3 (positive or negative)	Weak or No Correlation	Indicates a weak linear relationship or no linear relationship between the variables.

```
[ ]: corr = numerical_data.corr()

corr_sorted = corr["SalePrice"].sort_values(ascending=False)
corr_high = corr_sorted[abs(corr_sorted)>0.4].index.tolist()
corr_numVar = corr.loc[corr_high,corr_high]

sns.heatmap(corr_numVar, cmap="YlGnBu", annot=True)
plt.title("Correlation Coefficient of numerical variables and price")
plt.show()
```



Skewness

Skewness Range	Description
Near Zero (0)	Data is symmetric, with skewness between -0.5 and 0.5.
Moderate Skewness	Data is moderately asymmetric, with skewness between -1 to -0.5 or 0.5 to 1.
High Skewness	Data is highly skewed, with skewness greater than 1 or less than -1. Transformations needed.

```
[ ]: # Make a copy of numerical_data to avoid modifying the original dataset
numerical = numerical_data.drop(columns=["SalePrice"]).copy()
features = numerical.columns.tolist()
skewed_features = []

num_rows = 8
num_cols = 6
```

```

# Create subplots for visualization
fig, axes = plt.subplots(num_rows, num_cols, figsize=(25, 4*num_rows))
axes = axes.flatten()

# Iterate over each feature
for i, feature in enumerate(features):
    # Plot histogram with KDE (Kernel Density Estimate)
    sns.histplot(numerical[feature], kde=True, ax=axes[i])

    # Calculate skewness for the current feature
    skewness = skew(numerical[feature].dropna())

    # Identify and store features with skewness greater than 1
    if skewness > 1:
        skewed_features.append(feature)

    axes[i].set_title(f'{feature} - Skewness: {skewness:.2f}')
    axes[i].set_xlabel(feature)
    axes[i].set_ylabel('')

print("Skewed Features:", skewed_features)
for j in range(i + 1, len(axes)):
    fig.delaxes(axes[j])

plt.tight_layout()
plt.show()

```

```

Skewed Features: ['LotFrontage', 'LotArea', 'MasVnrArea', 'BsmtFinSF1',
'BsmtFinSF2', 'TotalBsmtSF', '1stFlrSF', 'LowQualFinSF', 'GrLivArea',
'BsmtHalfBath', 'KitchenAbvGr', 'WoodDeckSF', 'OpenPorchSF', 'EnclosedPorch',
'3SsnPorch', 'ScreenPorch', 'PoolArea', 'MiscVal']

```




Characteristics summary of numercial variable derived from EDA

Numercial Variable	Characteristics summary derived from EDA
LotFrontage	Skewed data, initially had outliers, strong correlation with SalePrice.
LotArea	Skewed data, initially had outliers, moderate correlation with SalePrice.
OverallQual	Strong positive correlation with SalePrice.
OverallCond	Weak negative correlation with SalePrice.
MasVnrArea	Skewed data, moderate correlation with SalePrice.
BsmtFinSF1	Skewed data, moderate correlation with SalePrice.
BsmtFinSF2	Skewed data, low correlation with SalePrice.
BsmtUnfSF	Low correlation with SalePrice.
TotalBsmtSF	Skewed data, initially had outliers, strong correlation with SalePrice.
1stFlrSF	Skewed data, moderate correlation with SalePrice.
2ndFlrSF	Low correlation with SalePrice.

Numercial Variable	Characteristics summary derived from EDA
LowQualFinSF	Low correlation with SalePrice.
GrLivArea	Skewed data, strong correlation with SalePrice.
BsmtFullBath	Low correlation with SalePrice.
BsmtHalfBath	Skewed data, low correlation with SalePrice.
FullBath	Moderate correlation with SalePrice.
HalfBath	Low correlation with SalePrice.
BedroomAbvGr	Low correlation with SalePrice.
KitchenAbvGr	Skewed data, weak negative correlation with SalePrice.
TotRmsAbvGrd	Moderate correlation with SalePrice.
Fireplaces	Moderate correlation with SalePrice.
GarageYrBlt	Low correlation with SalePrice.
GarageCars	Moderate correlation with SalePrice.
GarageArea	Moderate correlation with SalePrice.
WoodDeckSF	Skewed data, moderate correlation with SalePrice.
OpenPorchSF	Skewed data, moderate correlation with SalePrice.
EnclosedPorch	Low correlation with SalePrice.
3SsnPorch	Low correlation with SalePrice.
ScreenPorch	Low correlation with SalePrice.
PoolArea	Low correlation with SalePrice.
MiscVal	Skewed data, initially had outliers, low correlation with SalePrice.

2.4 Categorical Features

There are total 16 categorical features.

```
[ ]: categorical_data = df.select_dtypes(exclude="number")
categorical_list = categorical_data.columns.tolist()
print("Number of categorical variables " , len(categorical_list))
```

Number of categorical variables 16

```
[ ]: print(categorical_list)
```

```
['LandContour', 'BldgType', 'HouseStyle', 'RoofStyle', 'ExterQual', 'ExterCond',
'BsmtQual', 'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtFinType2',
'HeatingQC', 'KitchenQual', 'GarageFinish', 'GarageQual', 'GarageCond']
```

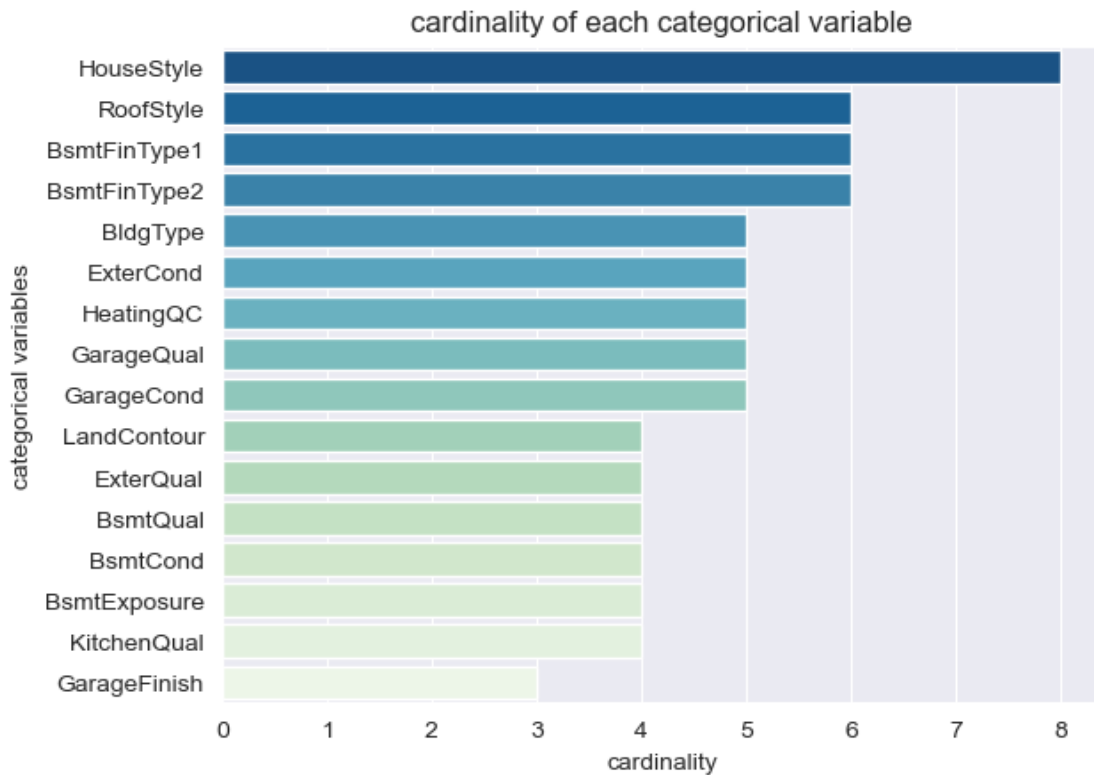
To have an initial feel of categorical variables 's characteristics. The cardinality of categorical variables is visualized. Categorical variables with high cardinality (i.e. HouseStyle, RoofStyle, BsmtFinType1, and BsmtFinType2) may have a significant impact on the analysis or modeling process.

```
[ ]: cardinality = {var: df[var].nunique() for var in categorical_list}
cardinality_df = pd.DataFrame(cardinality.items(), columns=["Variable",
↪ "Cardinality"])
cardinality_df=cardinality_df.sort_values(by="Cardinality",ascending=False)
```

```

sns.
    ↳ barplot(y=cardinality_df["Variable"],x=cardinality_df["Cardinality"],data=cardinality_df,
    ↳ palette="GnBu_r")
plt.ylabel("categorical variables")
plt.xlabel("cardinality")
plt.title("cardinality of each categorical variable")
plt.show()

```



3. Data Preprocessing

3.1 Missing Value

For numerical variables, given that the missing percentage for `LotFrontage` and `GarageYrBlt` is relatively high (17.7% and 5.5% respectively), creating missing value indicators is more suitable compared to other approaches such as mean imputation. Each of these variables are replaced by a new variable that works as a missing value indicator. If the value in the original variable is missing then the corresponding value in the new variable should be 1, otherwise it should be 0.

```

[ ]: numerical_missing = numerical_data.isna().sum().sort_values(ascending=False)
numerical_missing_percent = numerical_missing/len(numerical_data) * 100
numerical_missing_data = pd.DataFrame({
    "Percent": numerical_missing_percent,

```

```
numerical_missing_data.head(10)
"Count": numerical_missing})
```

```
[ ]:
      Percent  Count
LotFrontage   17.739726   259
GarageYrBlt    5.547945    81
MasVnrArea     0.547945     8
BedroomAbvGr   0.000000     0
MiscVal        0.000000     0
PoolArea       0.000000     0
ScreenPorch    0.000000     0
3SsnPorch      0.000000     0
EnclosedPorch  0.000000     0
OpenPorchSF    0.000000     0
```

```
[ ]: vars_with_missing = []
for var in df.select_dtypes(include="number").columns:
    if df[var].isnull().any():
        vars_with_missing.append(var)
for var in vars_with_missing:
    df[var+"_missing"] = df[var].isnull().astype(int)
    df.drop(columns=[var], inplace=True)

sum(df.select_dtypes(include="number").isna().sum())
```

```
[ ]: 0
```

For categorical variables, the proportion of missing values is relatively low (GarageFinish, GarageQual, and BsmtExposure range between 2.5% and 5.5%, this is considered low) and unlikely to skew the data. Mode imputation for categorical variables replaces missing values with the most frequent category.

```
[ ]: categorical_missing = categorical_data.isna().sum().sort_values(ascending=False)
categorical_missing_percent = categorical_missing/len(categorical_data) * 100
categorical_missing_data = pd.DataFrame({
    "Percent": categorical_missing_percent,
    "Count": categorical_missing})
categorical_missing_data.head(10)
```

```
[ ]:
      Percent  Count
GarageFinish   5.547945    81
GarageQual     5.547945    81
GarageCond     5.547945    81
BsmtExposure   2.602740    38
BsmtFinType2   2.602740    38
BsmtQual       2.534247    37
BsmtCond       2.534247    37
BsmtFinType1   2.534247    37
```

```
LandContour    0.000000    0
BldgType       0.000000    0
```

```
[ ]: for column in categorical_data.columns:
      mode_value = df[column].mode()[0]
      df[column].fillna(mode_value, inplace=True)
sum(df.select_dtypes(exclude="number").isna().sum())
```

```
[ ]: 0
```

3.2 Label Encoding

Many machine learning algorithms require numerical input. Label encoding is necessary to transform these categories into numerical values. All ordinal categorical variables are encoded into numbers based on the mean value of the target variable (`SalePrice`). The smaller value corresponds to the category that has the smaller mean house sale price. That is, the category that has the smallest mean house sale price can be replaced with 0, the next category with 1, and so on. Compared to one-hot encoding, which creates binary columns for each category, this approach reduces the dimensionality of the feature space.

```
[ ]: cat_vars = df.select_dtypes(exclude=['number']).columns

for var in cat_vars:
    # Calculate mean sale price for each category in the training set
    mean_sale_price = df.groupby(var)['SalePrice'].mean().sort_values()

    # Create a mapping dictionary for encoding
    encoding_map = {category: i for i, category in enumerate(mean_sale_price.
↳ index)}

    # Apply encoding to both training and testing datasets
    # Define the default value to assign to missing values in the test dataset
    # This value represents the encoding for missing categories that were not_
↳ present in the training dataset
    df["encoded_" + var] = df[var].map(encoding_map)
    df["encoded_" + var] = df[var].map(encoding_map)

    # Drop original categorical variables
    df.drop(columns=[var], inplace=True)
```

4. Feature Engineering

4.1 Split Data into Training and Testing Sets

Split the data into training and test sets, with 70% for training and 30% for testing. This step is done after data preprocessing but before transformation and standardization to prevent data leakage.

```
[ ]: from sklearn.model_selection import train_test_split


train_df, test_df = train_test_split(df, test_size=0.3, random_state=10)
```

4.2 Fixing skewness with Log Transformation


The target variable(**SalePrice**) is right skewed as shown before. Since normal distribution is crucial for linear regression, a log transformation is performed on skewed data to approximate a normal distribution. From distribution plot and QQ plot, the transformed data is approximately normally distributed.

```
[ ]: train_df['SalePrice'] = np.log(train_df['SalePrice'])
test_df['SalePrice'] = np.log(test_df['SalePrice'])

sns.displot(train_df["SalePrice"], kde=True)
plt.title("Distribution Plot")
plt.show()
```



```
[ ]: stats.probplot(train_df["SalePrice"], dist="norm", plot=plt)
plt.title("Probablity Plot")
plt.show()
```



4.3 Standardization and transformation

$$x'_i = \frac{x_i - \mu}{\sigma}$$

```
[ ]: from sklearn.preprocessing import PowerTransformer

y_train= train_df['SalePrice']
y_test= test_df['SalePrice']

X_train = train_df.drop('SalePrice', axis=1)
X_test = test_df.drop('SalePrice', axis=1)

# Standardize the features to have mean=0 and variance=1
# StandardScaler can help mitigate overflow issues by scaling the data
scaler = StandardScaler()
```

```
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
[ ]: from sklearn.preprocessing import PowerTransformer

# Apply the Yeo-Johnson transformation to the standardized features
# PowerTransformer with 'yeo-johnson' method handles non-positive data, but
  ↳ overflow may still occur
pt = PowerTransformer(method='yeo-johnson')
X_train_scaled = pd.DataFrame(pt.fit_transform(X_train_scaled), columns=X_train.
  ↳ columns)
X_test_scaled = pd.DataFrame(pt.transform(X_test_scaled), columns=X_test.
  ↳ columns)
```

```
[ ]: mean_before = X_train.mean()
std_before = X_train.std()

mean_after = X_train_scaled.mean()
std_after = X_train_scaled.std()

summary_df = pd.DataFrame(
    {'Feature': X_train.columns,
     'Mean Before': mean_before.values,
     'Std Before': std_before.values,
     'Mean After standardalization and transformation': mean_after,
     'Std After standardalization and transformation': std_after}
)

summary_df
```

```
[ ]:
```

	Feature	Mean Before	Std Before \
LotArea	LotArea	10681.630137	11305.457364
OverallQual	OverallQual	6.141879	1.372897
OverallCond	OverallCond	5.562622	1.087500
BsmtFinSF1	BsmtFinSF1	447.314090	468.198594
BsmtFinSF2	BsmtFinSF2	40.860078	150.289590
BsmtUnfSF	BsmtUnfSF	575.344423	442.929842
TotalBsmtSF	TotalBsmtSF	1063.518591	443.412565
1stFlrSF	1stFlrSF	1167.607632	393.824408
2ndFlrSF	2ndFlrSF	347.983366	436.440810
LowQualFinSF	LowQualFinSF	7.010763	53.596059
GrLivArea	GrLivArea	1522.601761	533.663862
BsmtFullBath	BsmtFullBath	0.422701	0.517465
BsmtHalfBath	BsmtHalfBath	0.053816	0.230062
FullBath	FullBath	1.573386	0.549237
HalfBath	HalfBath	0.388454	0.505391
BedroomAbvGr	BedroomAbvGr	2.873777	0.816527

KitchenAbvGr	KitchenAbvGr	1.044031	0.205265
TotRmsAbvGrd	TotRmsAbvGrd	6.526419	1.598477
Fireplaces	Fireplaces	0.615460	0.644327
GarageCars	GarageCars	1.775930	0.748284
GarageArea	GarageArea	478.349315	215.881679
WoodDeckSF	WoodDeckSF	95.712329	125.395412
OpenPorchSF	OpenPorchSF	47.555773	66.206621
EnclosedPorch	EnclosedPorch	21.959883	61.819216
3SsnPorch	3SsnPorch	3.849315	32.068988
ScreenPorch	ScreenPorch	13.813112	51.399277
PoolArea	PoolArea	2.799413	40.430832
MiscVal	MiscVal	54.387476	585.607282
LotFrontage_missing	LotFrontage_missing	0.173190	0.378597
MasVnrArea_missing	MasVnrArea_missing	0.004892	0.069808
GarageYrBlt_missing	GarageYrBlt_missing	0.053816	0.225765
encoded_LandContour	encoded_LandContour	1.045988	0.439079
encoded_BldgType	encoded_BldgType	3.679061	0.851761
encoded_HouseStyle	encoded_HouseStyle	4.800391	1.385980
encoded_RoofStyle	encoded_RoofStyle	1.587084	1.198166
encoded_ExterQual	encoded_ExterQual	1.409980	0.574702
encoded_ExterCond	encoded_ExterCond	2.868885	0.393919
encoded_BsmtQual	encoded_BsmtQual	1.568493	0.686162
encoded_BsmtCond	encoded_BsmtCond	2.007828	0.269103
encoded_BsmtExposure	encoded_BsmtExposure	0.651663	1.026577
encoded_BsmtFinType1	encoded_BsmtFinType1	3.386497	1.624474
encoded_BsmtFinType2	encoded_BsmtFinType2	3.760274	0.850343
encoded_HeatingQC	encoded_HeatingQC	3.164384	0.951510
encoded_KitchenQual	encoded_KitchenQual	1.524462	0.670410
encoded_GarageFinish	encoded_GarageFinish	0.789628	0.813554
encoded_GarageQual	encoded_GarageQual	1.980431	0.241625
encoded_GarageCond	encoded_GarageCond	3.899217	0.542620

Mean After standardalization and transformation \

LotArea	7.332686e-18
OverallQual	4.323569e-17
OverallCond	-2.169932e-16
BsmtFinSF1	3.454510e-17
BsmtFinSF2	-3.653307e-16
BsmtUnfSF	1.151503e-17
TotalBsmtSF	2.707662e-17
1stFlrSF	-7.020368e-18
2ndFlrSF	2.752745e-16
LowQualFinSF	8.127875e-16
GrLivArea	1.596896e-17
BsmtFullBath	-2.628904e-17
BsmtHalfBath	1.565854e-15
FullBath	-7.093695e-17

HalfBath	-3.354568e-16
BedroomAbvGr	-1.295441e-16
KitchenAbvGr	1.645781e-17
TotRmsAbvGrd	-1.412221e-17
Fireplaces	8.147429e-17
GarageCars	-7.843259e-17
GarageArea	4.562560e-18
WoodDeckSF	4.169311e-16
OpenPorchSF	6.452764e-17
EnclosedPorch	8.505916e-17
3SsnPorch	2.850731e-15
ScreenPorch	1.735946e-16
PoolArea	9.742153e-16
MiscVal	1.040101e-15
LotFrontage_missing	-2.400776e-17
MasVnrArea_missing	-2.096361e-15
GarageYrBlt_missing	1.124454e-15
encoded_LandContour	-1.198534e-16
encoded_BldgType	3.574006e-17
encoded_HouseStyle	-8.820950e-17
encoded_RoofStyle	-1.435577e-16
encoded_ExterQual	-1.218855e-16
encoded_ExterCond	-5.040543e-17
encoded_BsmtQual	-8.397284e-17
encoded_BsmtCond	1.119117e-16
encoded_BsmtExposure	3.091678e-16
encoded_BsmtFinType1	-1.306304e-17
encoded_BsmtFinType2	-7.089893e-16
encoded_HeatingQC	-9.364112e-17
encoded_KitchenQual	-9.342386e-18
encoded_GarageFinish	-2.096605e-16
encoded_GarageQual	-3.087740e-16
encoded_GarageCond	2.363298e-16

Std After standardalization and transformation

LotArea	1.00049
OverallQual	1.00049
OverallCond	1.00049
BsmtFinSF1	1.00049
BsmtFinSF2	1.00049
BsmtUnfSF	1.00049
TotalBsmtSF	1.00049
1stFlrSF	1.00049
2ndFlrSF	1.00049
LowQualFinSF	1.00049
GrLivArea	1.00049
BsmtFullBath	1.00049

BsmtHalfBath	1.00049
FullBath	1.00049
HalfBath	1.00049
BedroomAbvGr	1.00049
KitchenAbvGr	1.00049
TotRmsAbvGrd	1.00049
Fireplaces	1.00049
GarageCars	1.00049
GarageArea	1.00049
WoodDeckSF	1.00049
OpenPorchSF	1.00049
EnclosedPorch	1.00049
3SsnPorch	1.00049
ScreenPorch	1.00049
PoolArea	1.00049
MiscVal	1.00049
LotFrontage_missing	1.00049
MasVnrArea_missing	1.00049
GarageYrBlt_missing	1.00049
encoded_LandContour	1.00049
encoded_BldgType	1.00049
encoded_HouseStyle	1.00049
encoded_RoofStyle	1.00049
encoded_ExterQual	1.00049
encoded_ExterCond	1.00049
encoded_BsmtQual	1.00049
encoded_BsmtCond	1.00049
encoded_BsmtExposure	1.00049
encoded_BsmtFinType1	1.00049
encoded_BsmtFinType2	1.00049
encoded_HeatingQC	1.00049
encoded_KitchenQual	1.00049
encoded_GarageFinish	1.00049
encoded_GarageQual	1.00049
encoded_GarageCond	1.00049

```
[ ]: # Box plot comparing original, standardized, and transformed distributions
plt.figure(figsize=(16, 24))
sns.boxplot(data=pd.DataFrame(X_train_scaled, columns=X_train.columns), orient='h', palette='Set3', showfliers=False)
plt.title('Distribution Comparison before Standardization and Yeo-Johnson Transformation')
plt.xticks()
plt.show()
```

